

Arshan Bhanage

San Jose, CA | (408) 401-6724 | arshan.bhanage@sjsu.edu | [linkedin.com/in/arshanbhanage](https://www.linkedin.com/in/arshanbhanage)

EDUCATION

M.S., Computer Software Engineering

Expected May 2027

San Jose State University, San Jose, CA

Coursework: Advanced Data Mining, Enterprise Software Platforms, Data Structures & Algorithms, Artificial Intelligence

TECHNICAL SKILLS

Languages: Java, C++, Python, Ruby, GoLang, TypeScript (basic)

Backend: Spring Boot (MVC), Object Oriented Design, Node.js, REST APIs, Microservices

AI & Data: TensorFlow, LLM, RAG, Agentic AI, Kafka, PostgreSQL, Vector Databases

Misc: AWS, Git (CI/CD), SDLC, Docker

Core Competencies: Growth Mindset, Customer Obsession, Complex Problem Solving, Agile Collaboration

EXPERIENCE

Software Engineer, IBM, Mumbai, India

Jun 2024 - Jun 2025

- Built modular Java/Spring Boot microservices for Union Bank of India's POS onboarding platform, supporting 10k+ merchants and enabling 30% faster POS activation with 20% fewer application rejections.
- Executed core flows—user registration, application submission, and status transitions—and emitted settlement events to Apache Kafka to drive downstream processing.
- Reduced friction in onboarding by launching actionable feedback and re-application flows plus implemented status tracking/notifications scaled to 10,000+ merchants.

Software Engineer Intern, Vmayo Technologies, India

May 2023 - Oct 2023

- Constructed on a complex Ruby on Rails open-source codebase (Postal) to deliver a scalable, multi-tenant, white-label SaaS for resellers—launched a new reseller model with 90% client retention.
- Extended React components and Rails services to support per-tenant overrides for branding, custom domains, and IP-pool management, enabling reseller-specific configuration at scale.
- Designed PostgreSQL schema extensions and RESTful APIs to provision 15+ resellers and process 200k+ emails/day; owned design-to-deployment.
- Collaborated via code reviews and Agile sprint planning, navigating a large, complex codebase.

KEY PROJECTS

Driftwise - Real-Time Contract Compliance, Financial Agent Hackathon Finalist

Oct 2025 - Present

- Established and scaled AI infrastructure for continuous compliance monitoring of financial contracts (loans, insurance, filings) with Python + LangGraph + Claude Sonnet 4.5 agent plans; achieved evidence-linked decisions and sub-second violation detection and sustained 1k+ events/min (99.9% pipeline uptime).
- Developed a Policy Compiler and Validation Engine, translated legal clauses into executable rules and streamed results via FastAPI SSE, cutting rule creation time by 90% and facilitating instant compliance visibility.
- Accomplished unstructured document parsing with LandingAI ADE pipeline (Parse, Extract), facilitating traceable, visual-grounded outputs; median extraction <15s/doc at scale and throughput scaling to 100+ pages/min with parallelism.

Parallel Log Analyzer, Personal Project

Nov 2025 - Present

- Developed a high-performance, multi-threaded command-line utility in C++ to parse and aggregate statistics from large-scale server log files for an Advanced Algorithms course project.
- Implemented a parallel processing pipeline using std::thread and thread-safe queues, catalyzing tool to process 1GB of log data in under 10 seconds—a 40x improvement over a single-threaded Python baseline.
- Optimized memory usage by over 70% through efficient use of STL containers (e.g., unordered map) and smart pointers, allowing analysis of multi-gigabyte files on standard hardware.

TraceQA - Enterprise Requirements AI, Code Rabbit Hackathon Top 6/180

Sep 2025 - Oct 2025

- Installed an on-prem, privacy-preserving LLM system (banking) with Python & FastAPI microservices for ingest, chunk/embedding, and RAG retrieval on Pinecone; achieved document-query p95 <900ms with retrieval and lightweight re-ranking.
- Addressed requirement versioning, semantic diffing, and impact analysis with automated test generation; produced 500+ tests and raised coverage from 62% to 92%.
- Tuned RAG (chunking, retrieval, prompt scaffolds) lifted MRR by 28% on internal evaluations; background jobs and dashboards shortened feedback loops from hours to minutes.